

# STAT S 670 - Exploratory Data Analysis - Midterm

Ganesh Nagarajan, [gnagaraj@indiana.edu](mailto:gnagaraj@indiana.edu)

November 3, 2015

1. (a) Calculate 5 Number summary for wt2

```
wt2<-c(143,-184,182,-110,1017,986,1010,1001,-111,-60,-151,-111,1024,1031,1028)
summary(wt2)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -184.0  -110.5   182.0   446.3  1014.0  1031.0
```

- (b) Stem leaf plot

According to Tukey, the range of the dataset determines the lines on the stem leaf plot. Range can be given as follows,

```
library(aplpack)
```

```
## Loading required package: tcltk
```

```
range.wt2<-(range(wt2)[2]-range(wt2)[1])/length(wt2)
range.wt2
```

```
## [1] 81
```

```
stem.leaf(wt2,100)
```

```
## 1 | 2: represents 1200
## leaf unit: 100
##           n: 15
##      6   -0* | 111110
##    (2)   0* | 11
##           t |
##           f |
##           s |
##      7    0. | 9
##      6    1* | 000000
```

Another way of representing this is using the stem command,

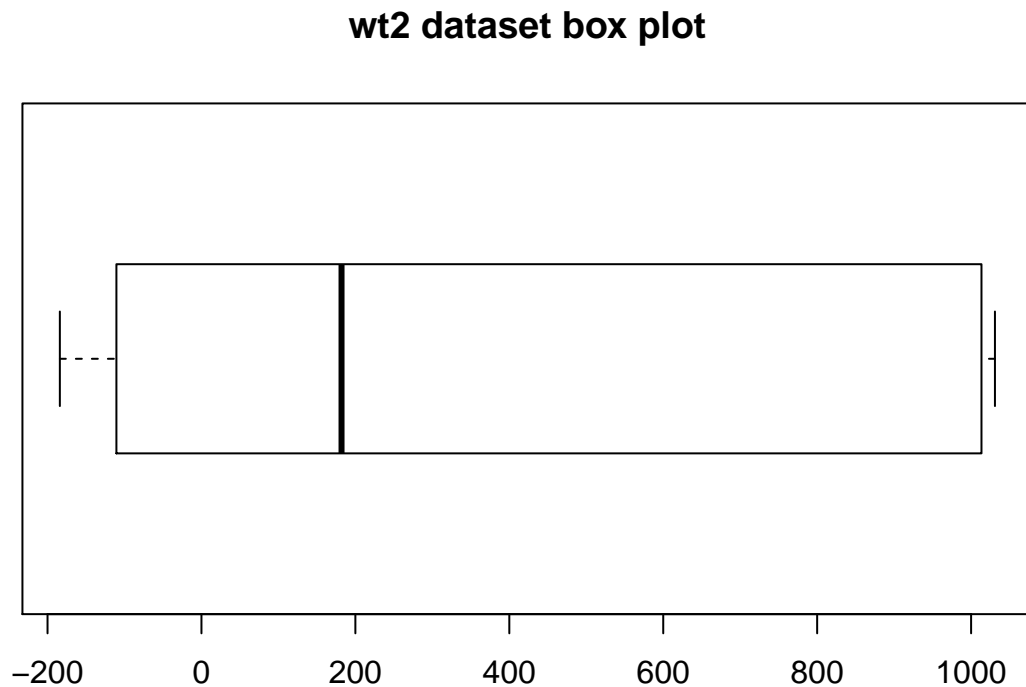
```
stem(wt2,2)
```

```
##
## The decimal point is 2 digit(s) to the right of the |
##
## -0 | 851116
```

```
##    0 | 48
##    2 |
##    4 |
##    6 |
##    8 | 9
##   10 | 012233
```

Below is the box plot of wt2 and it can be seen that there are no outliers.

```
boxplot(wt2, horizontal = TRUE, main="wt2 dataset box plot")
```



## 2. “Outsides values in a box plot”


The average outside values in a batch can be given by the formula  $0.4 + 0.07n$  where  $n$  is the number of elements in the batch. considering 5000 elements, outliers are as follows,

```
noOutliers<-function(x){0.4+(0.007*x)}
noOutliers(5000)
```

```
## [1] 35.4
```

Lets check the correctness of above predicted number. Lets generate 5000 random values and check the outliers.

```
set.seed(3)
boxplot(rnorm(5000,mean=5,sd=0.5),horizontal = TRUE,main="Simulation of Random 5000 samples")
```



Above graph has about 30 outliers which is right about the estimate we have done.

### 3. PENDING

4. Smoothers in general tries to identify the patterns in the data. They remove noise signals and try to replicate the trends of the actual data. Linear smoothers are smoothers which transform the data linearly. This is usually a convolution function.

Examples of Linear Smoothers,

1. Running Mean
2. Cubic Spline Smoothers
3. Lowess Smoother

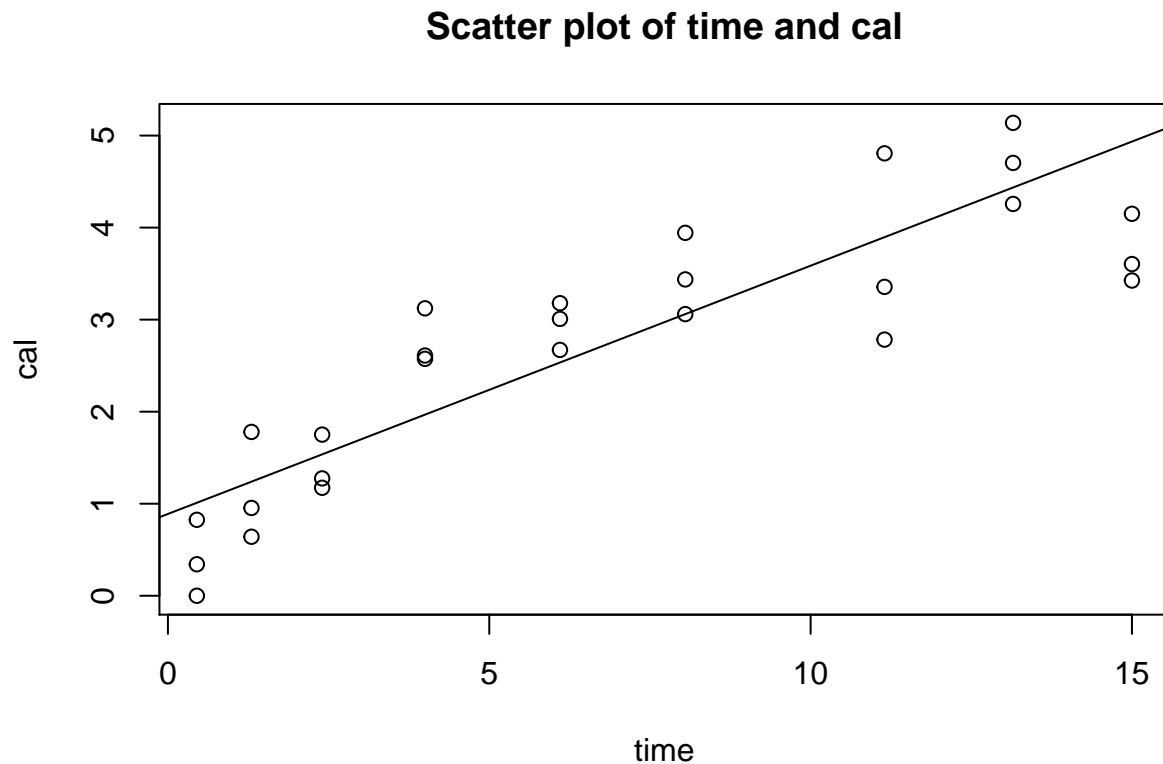
Disadvantages of Linear Smoother,

1. Over fits the data, attempts to fit every point available
2. Since it attempts to fit every point, they are affected by outliers
3. Some linear smoothers like LOWESS filters require fairly dense points for smoothing.

Advantage and Disadvantage of Non linear smoothers

1. Since Non linear smoothers doesn't have linear constraints, they are quite flexible
2. Since they involve medians, they are more robust than the other methods and are less affected by the outliers
3. Requires more computation than linear smoothers. Usually require denser data.
4. (a) Initial RR line on the plot.

```
source("myplotfit.r")
source("rrline.r")
x<-c(0.450,0.45,0.450,1.300,1.300,1.300,2.400,2.400,2.400,4.000,4.00,4.000,6.100,6.100,6.100,8.05,8.050
y<-c(0.342,0.00,0.825,1.780,0.954,0.641,1.751,1.275,1.173,3.123,2.61,2.574,3.179,3.008,2.671,3.06,3.943
#Run one iteration of RR
rrline.x<-rrline1(x,y)
plot(x,y,main="Scatter plot of time and cal",xlab = "time",ylab = "cal")
abline(rrline.x$a,rrline.x$b)
```




```
paste("Intercept : ",rrline.x$a," Slope : ",rrline.x$b)
```

```
## [1] "Intercept : 0.888877637130801 Slope : 0.269704641350211"
```

(b) PENDING

(c) The Y values are as it is, however X axis is transformed. The Transformed used is square root.


```
sqrtx<-sqrt(x)
plot(sqrtx,y,main="Transformed plot",xlab="log(x)",ylab="y")
```



(d) RR on panel b

```
sqrtx<-sqrt(x)
plot(sqrtx,y,main="Transformed plot",xlab="log(x)",ylab="y")
rrline.sqrtx<-rrline1(sqrtx,y)
abline(rrline.sqrtx$a,rrline.sqrtx$b,lty=2,col="red")
```

## Transformed plot



```
paste("Intercept : ",rrline.sqrtrx$a," Slope : ",rrline.sqrtrx$b)
```

```
## [1] "Intercept : -0.404739545752101 Slope : 1.28553862281732"
```

(e) A general observation is that the residual plot of the panel b has more random points than the residual plot of panel a. More random the residual plot, better the fit model.

```
rrline.x$sumres
```

```
## [1] 16.38401
```

```
rrline.sqrtrx$sumres
```

```
## [1] 13.13373
```

Also an observation is that sum of residuals of plot b is lesser than plot a, ie the model fits better than the plot 1.

(f) RR after two iterations.

```

source("myplotfit.r")
source("rrline.r")
x<-c(0.450,0.45,0.450,1.300,1.300,1.300,2.400,2.400,2.400,4.000,4.00,
      4.000,6.100,6.100,6.100,8.05,8.050,8.050,11.150,11.150,11.150,
      13.150,13.150,13.150,15.000,15.00,15.000)
y<-c(0.342,0.00,0.825,1.780,0.954,0.641,1.751,1.275,1.173,3.123,2.61,
      2.574,3.179,3.008,2.671,3.06,3.943,3.437,4.807,3.356,2.783,5.138,
      4.703,4.257,3.604,4.15,3.425)

intercept <- 0.256
slope <- -0.124
r.new <- slope*x+intercept
#Iteration 2
rr1<-rrline2(x,r.new)
paste("Intercept : ",rr1$a," Slope : ",rr1$b)

## [1] "Intercept : 0.256 Slope : -0.124"

```

#### 6. (a) Median Polish

```

source("myplotfit.r")
source("rrline.r")
red<-c(5,6,3,11,10)
white<-c(14,10,6,12,21)
pink<-c(16,24,15,26,32)
mat<-rbind(red,white,pink)
colnames(mat)<-c("A1","A2","A3","A4","A5")
res<-medpolish(mat)

```

```

## 1: 28
## Final: 28

```

```
print(res)
```

```

##
## Median Polish Results (Dataset: "mat")
##
## Overall: 12
##
## Row Effects:
##   red white  pink
##   -6     0    12
##
## Column Effects:
## A1 A2 A3 A4 A5
## -1  0 -6  2  8
##
## Residuals:
##      A1 A2 A3 A4 A5
## red    0  0  3  3 -4
## white  3 -2  0 -2  1
## pink  -7  0 -3  0  0

```

The row effects suggest that Pink area has the maximum visit (12) and column effects suggests that Area 5 (8) have highest visitors.


- (b) Quality of the fit can be given by Analog R2. Its curious that the overall effects (12) is same as the median (12).

```
Analog_R_Square<- 1-((sum(abs(res$residuals))) / (sum(abs(mat-res$overall))))
paste("Analog R2 : ",Analog_R_Square)
```

```
## [1] "Analog R2 : 0.711340206185567"
```

- (c) Diagnostic Plot

```
diag.MP(res)
```



```
##      a      b  |res|
##  1  0  1.00000 31.50000
##  2  0 -1.16667 31.91667
##  3  0  1.66667 43.25000
##  4  0 -1.50000 28.00000
##  5  0  1.00000 31.50000
##  6  0 -1.16667 31.91667
##  7  0  1.66667 43.25000
##  8  0 -1.50000 28.00000
```



```
## 9 0 1.00000 31.50000
## 10 0 -1.16667 31.91667
## 0 -0.16667 31.91667
```

Diagnostic plot tells us that how good the model is. If the residuals are along 0, Then the model is good.

Assuming the model is a perfect fit, a perfect fit matrix is calculated as below.

$$\text{Fit } Y_{ij} = m + a_i + b_j + (a_i b_j) / m.$$

	A <sub>1</sub>	A <sub>2</sub>	A <sub>3</sub>	A <sub>4</sub>	A <sub>5</sub>
Red	5.5	6	3	7	10.25
White	11	12	6	14	20
Pink	22	24	12	28	40

Actual Residuals:


	A <sub>1</sub>	A <sub>2</sub>	A <sub>3</sub>	A <sub>4</sub>	A <sub>5</sub>
Red	0	0	3	3	-4
White	3	-2	0	-2	1
Pink	-7	0	-3	0	0

plot (0, 5.5) (0, 6) (3, 3) (3, 7) (-4, 10.25) etc.

A full fit,  $y = m + a_i + b_j + (a_i b_j) / m$  and the residuals are plotted as pairs as explained above.

(d) Forget it plot and Symbol Plot

```
forgetitplot(res)
```




An interpretation of the plot is that A5 is farthest away from any immediate areas and Pink is the farthest away from other immediate points. Hence these these points have the maximum values of incoming insects.

7 (a) PENDING

(b) K Smooth for various lambda values for normal kernel.

```
mu<-read.csv("smoothing.csv",header = FALSE)
x<-mu$V1
y<-mu$V2
mu = function(t){t + 0.5 *exp(-50*(t-0.5)^2)}
curve(mu(x),0,1,col="red",main="Kernel Smoothing with normal kernel")
points(x,y)
t=ksmooth(x,y,kernel="normal",bandwidth=0.5)
lines(t$x,t$y,col="violet",lty=2)
text(t$x[1],t$y[1],"Bandwidth = 0.5",col="violet")
t=ksmooth(x,y,kernel="normal",bandwidth=0.3)
lines(t$x,t$y,col="orange",lty=2)
text(t$x[20],t$y[20],"Bandwidth = 0.3",col="orange")
t=ksmooth(x,y,kernel="normal",bandwidth=0.25)
lines(t$x,t$y,col="blue")
text(t$x[30],t$y[30],"Bandwidth = 0.25",col="blue")
t=ksmooth(x,y,kernel="normal",bandwidth=0.15)
lines(t$x,t$y,lty=5)
text(t$x[40],t$y[40],"Bandwidth = 0.15")
```


## Kernel Smoothing with normal kernel



K Smooth for various lambda values for box kernel

```
mu<-read.csv("smoothing.csv",header = FALSE)
x<-mu$V1
y<-mu$V2
mu = function(t){t + 0.5 *exp(-50*(t-0.5)^2)}
curve(mu(x),0,1,col="red",main="Kernel Smoothing with box kernel")
points(x,y)
t=ksmooth(x,y,kernel="box",bandwidth=0.5)
lines(t$x,t$y,col="violet",lty=2)
text(t$x[1],t$y[1],"Bandwidth = 0.5",col="violet")
t=ksmooth(x,y,kernel="box",bandwidth=0.3)
lines(t$x,t$y,col="orange",lty=2)
text(t$x[20],t$y[20],"Bandwidth = 0.3",col="orange")
t=ksmooth(x,y,kernel="box",bandwidth=0.25)
lines(t$x,t$y,col="blue")
text(t$x[30],t$y[30],"Bandwidth = 0.25",col="blue")
t=ksmooth(x,y,kernel="box",bandwidth=0.15)
lines(t$x,t$y,lty=5)
text(t$x[40],t$y[40],"Bandwidth = 0.15")
```

## Kernel Smoothing with box kernel



(c) Function for calculating cross validation

```
cv <- function(x, y, lambda)
{
  n <- length(x)
  cv.total <- numeric(n)
  for(i in 1:n)
  {
    fit = ksmooth(x[-i], y[-i], kernel = "normal", bandwidth = lambda, n.points = n)
    cv.total[i] = (y[i] - fit$y[i])^2
  }
  mean.cv <- mean(cv.total)
  return(mean.cv)
}
```

(d) Plot for CV values vs lambda


```
x <- seq(0.01, 0.99, by = 0.02)
y <- c(-.0937, .0247, .1856, .1620, -.0316, .1442, .0993, .3823, -.0624, .3262, .1271, -.4158, .0975, -

lambda = seq(0.01, 1, by = 0.01)
cv.plot = c()
for(i in 1:length(lambda))
{
```

```

cv.temp = cv(x, y, lambda[i])
cv.plot = c(cv.plot, cv.temp)
}
plot(lambda, cv.plot,main="Lambda Vs CV Values")

```



```

lmin = lambda[which.min(cv.plot)]
paste("Lambda Minimum is : ",lmin)

```

```
## [1] "Lambda Minimum is : 0.14"
```

(e) Smoothing Splines

```

x = seq(0.01, 0.99, by = 0.02)
y = c(-.0937, .0247, .1856, .1620, -.0316, .1442, .0993, .3823, -.0624, .3262, .1271, -.4158, .0975, -.0937, .1442, .0993, .3823, -.0624, .3262, .1271, -.4158, .0975, -.0937)
mu = function(t){t + 0.5 *exp(-50*(t-0.5)^2)}
curve(mu(x),0,1,col="red",main="Kernel Smoothing Smoothing Spline")
points(x,y)
t=smooth.spline(x,y,spar = 0.5)
lines(t$x,t$y,col="violet",lty=2)
text(t$x[1],t$y[1], "Bandwidth = 0.5",col="violet")
t=smooth.spline(x,y,spar = 0.3)
lines(t$x,t$y,col="orange",lty=2)
text(t$x[20],t$y[20], "Bandwidth = 0.3",col="orange")
t=smooth.spline(x,y,spar = 0.8)


```

```

lines(t$x,t$y,col="blue")
text(t$x[30],t$y[30],"Bandwidth = 0.25",col="blue")
t=smooth.spline(x,y,spar = 0.2)
lines(t$x,t$y,lty=5)
text(t$x[40],t$y[40],"Bandwidth = 0.15")

```

## Kernel Smoothing Smoothing Spline



From the above graph, it can be seen that as the smoothing parameter decreases, the smoother tries to capture all the points, hence is more distorted

```

cv.spline <- function(x, y, lambda)
{
  n <- length(x)
  cv.total <- numeric(n)
  for(i in 1:n)
  {
    fit = smooth.spline(x[-i], y[-i],spar = lambda)
    cv.total[i] = fit$cv.crit
  }
  mean.cv<-mean(cv.total,na.rm=TRUE)
  return(mean.cv)
}

cv.spline.gcv <- function(x, y, lambda)
{
  n <- length(x)

```

```

cv.total <- numeric(n)
for(i in 1:n)
{
  fit = smooth.spline(x[-i], y[-i],spar = lambda,cv=FALSE)
  cv.total[i] = fit$cv.crit
}
mean.cv<-mean(cv.total,na.rm=TRUE)
return(mean.cv)
}

```

```

x = seq(0.01, 0.99, by = 0.02)
y = c(-.0937, .0247, .1856, .1620, -.0316, .1442, .0993, .3823, -.0624, .3262, .1271, -.4158, .0975, -.

lambda = seq(0.01, 1, by = 0.01)
cv.plot.spline = c()
cv.plot.spline.gcv = c()
for(i in 1:length(lambda))
{
  cv.temp = cv.spline(x, y, lambda[i])
  cv.plot.spline = c(cv.plot.spline, cv.temp)
}
lmin = lambda[which.min(cv.plot.spline)]
paste("Lambda Minimum is : ",lmin)

```

```
## [1] "Lambda Minimum is : 0.65"
```

```

for(i in 1:length(lambda))
{
  cv.temp = cv.spline.gcv(x, y, lambda[i])
  cv.plot.spline.gcv = c(cv.plot.spline.gcv, cv.temp)
}
lmin = lambda[which.min(cv.plot.spline.gcv)]
paste("Lambda Minimum for GCV is : ",lmin)

```


```
## [1] "Lambda Minimum for GCV is : 0.65"
```

```

plot(lambda, cv.plot.spline,main="Lambda Vs CV/GCV Values",col="red")
lines(lambda,cv.plot.spline.gcv,col="blue")

```

**Lambda Vs CV/GCV Values**



GCV values are